

```
import pandas as pd
```

```
df=pd.read_csv("sales_data_sample.csv",encoding = 'Latin-1')  
## if we don't add "",encoding = 'Latin-1'"" then we get error
```

```
df.head()
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	
SALES \					
0	10107	30	95.70	2	2871.00
1	10121	34	81.35	5	2765.90
2	10134	41	94.74	2	3884.34
3	10145	45	83.26	6	3746.70
4	10159	49	100.00	14	5205.27

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	\
0	2/24/2003 0:00	Shipped	1	2	2003	...	
1	5/7/2003 0:00	Shipped	2	5	2003	...	
2	7/1/2003 0:00	Shipped	3	7	2003	...	
3	8/25/2003 0:00	Shipped	3	8	2003	...	
4	10/10/2003 0:00	Shipped	4	10	2003	...	

	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME
DEALSIZE					
0	10022	USA	NaN	Yu	Kwai
Small					
1	51100	France	EMEA	Henriot	Paul
Small					
2	75508	France	EMEA	Da Cunha	Daniel
Medium					
3	90003	USA	NaN	Young	Julie
Medium					
4	NaN	USA	NaN	Brown	Julie
Medium					

```
[5 rows x 25 columns]
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ORDERNUMBER           2823 non-null   int64
1   QUANTITYORDERED       2823 non-null   int64
2   PRICEEACH             2823 non-null   float64
3   ORDERLINENUMBER       2823 non-null   int64
4   SALES                 2823 non-null   float64
5   ORDERDATE             2823 non-null   object
6   STATUS                2823 non-null   object
7   QTR_ID               2823 non-null   int64
8   MONTH_ID             2823 non-null   int64
9   YEAR_ID              2823 non-null   int64
10  PRODUCTLINE          2823 non-null   object
11  MSRP                 2823 non-null   int64
12  PRODUCTCODE          2823 non-null   object
13  CUSTOMERNAME         2823 non-null   object
14  PHONE                2823 non-null   object
15  ADDRESSLINE1         2823 non-null   object
16  ADDRESSLINE2         302 non-null    object
17  CITY                 2823 non-null   object
18  STATE                1337 non-null   object
19  POSTALCODE           2747 non-null   object
20  COUNTRY              2823 non-null   object
21  TERRITORY            1749 non-null   object
22  CONTACTLASTNAME      2823 non-null   object
23  CONTACTFIRSTNAME     2823 non-null   object
24  DEALSIZE             2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

```
df.isnull().sum()
```

ORDERNUMBER	0
QUANTITYORDERED	0
PRICEEACH	0
ORDERLINENUMBER	0
SALES	0
ORDERDATE	0
STATUS	0
QTR_ID	0
MONTH_ID	0
YEAR_ID	0
PRODUCTLINE	0
MSRP	0
PRODUCTCODE	0
CUSTOMERNAME	0
PHONE	0

```

ADDRESSLINE1      0
ADDRESSLINE2    2521
CITY              0
STATE            1486
POSTALCODE        76
COUNTRY           0
TERRITORY        1074
CONTACTLASTNAME   0
CONTACTFIRSTNAME  0
DEALSIZE          0
dtype: int64

```

```

## so only two columns are important of the dataset i.e.
QuantityOrdered and PriceEach other are irrelevant
data=df[['QUANTITYORDERED', 'PRICEEACH']]

```

```
data.head(4)
```

	QUANTITYORDERED	PRICEEACH
0	30	95.70
1	34	81.35
2	41	94.74
3	45	83.26

```

## Do normalization of the data
from sklearn.preprocessing import StandardScaler
# make object of it
scaler=StandardScaler()
normalized_data=scaler.fit_transform(data)
print(normalized_data)

```

```

[[-0.52289086  0.5969775 ]
 [-0.11220131 -0.11445035]
 [ 0.60650538  0.54938372]
 ...
 [ 0.81185016  0.81015797]
 [-0.11220131 -1.06186404]
 [ 1.2225397  -0.89925195]]

```

```

## Using elbow method , determine the best value of k
# wcss= within cluster sum of squares . It's a measure of how close
data points are to the centroid of their cluster

```

```
from sklearn.cluster import KMeans
```

```

wcss=[]
for i in range(1,16):
    k_means=KMeans(n_clusters=i,init='k-means+
+',max_iter=300,n_init=10,random_state=10)
    # Here
    # n_clusters specifies the number of clusters you want the

```

algorithm to find in your data.

init determines the method for initializing the positions of the cluster centers (centroids). 'k-means++' is the default and recommended method.

max_iter=300 Sets the maximum number of iterations the algorithm will run for a single initialization.

n_init=10 Specifies the number of times the KMeans algorithm will run with different centroid seeds.

```
k_means.fit(normalized_data)
wcss.append(k_means.inertia_)
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
```

```
warnings.warn(
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
```

```
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP_NUM_THREADS=12.  
warnings.warn(  
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable
```


```
OMP_NUM_THREADS=12.  
warnings.warn(  

```

```
## Elbow Graph
```

```
import matplotlib.pyplot as plt  
plt.plot(range(1,16),wcss)  
plt.xlabel("K value")  
plt.ylabel("WCSS")  
plt.title("Elbow Curve Graph")
```

```
Text(0.5, 1.0, 'Elbow Curve Graph')
```



```
# from graph we can see that for k=4 is the optimal value , so train  
the model
```

```
k_means=KMeans(n_clusters=4,init='k-means+  
+',max_iter=300,n_init=10,random_state=10)  
clusters=k_means.fit_predict(normalized_data)
```

```
# The fit_predict method combines the operations of fitting the model  
and predicting
```

```
C:\Users\Ashvini Mahajan\Anaconda\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1382: UserWarning: KMeans is known to have a memory leak on  
Windows with MKL, when there are less chunks than available threads.
```

You can avoid it by setting the environment variable
OMP_NUM_THREADS=12.
warnings.warn()

```
## Visualization of the clusters
plt.scatter(normalized_data[:, 0], normalized_data[:, 1], c=clusters,
            cmap='viridis')
# x axis numbers =normalized_data[:, 0]
# y axis numbers =normalized_data[:, 1]
# The c parameter specifies the color of the markers (data points) in
# the scatter plot
# The cmap parameter stands for "colormap." It defines the colormap
# used to map numerical data to colors.
plt.xlabel('QUANTITYORDERED')
plt.ylabel('PRICEEACH')
plt.title('K-Means Clustering')
plt.show()
```

